


# YOLO v8

Reconocimiento de objetos

Máster en Robótica y Sistemas Inteligentes




### Classification



Capacitor

### Object Detection



Capacitor, Resistor, Transformer, Connector,  
Inductor, Polyester Capacitor


### Segmentation



Capacitor, Resistor, Transformer, Connector,  
Inductor, Polyester Capacitor

# YOLO v8

## Installation



***Prerequisites: Python>=3.7 and PyTorch>=1.7***

Create conda environment

```
conda create --name reob_env python=3.10
```

Conda activate environment

```
conda activate reob_env
```

Conda install Pytorch (<https://pytorch.org/get-started/locally/>)

```
conda install pytorch torchvision torchaudio pytorch-cuda=11.7 -c pytorch -c nvidia
```

Install yolov8

```
pip install ultralytics
```


Install jupyter notebook

```
pip install notebook
```



# YOLO v8

## Inference



```
In [1]: 1 from ultralytics import YOLO
        2 import cv2
        3 from matplotlib import pyplot as plt
        4 import numpy as np
        5 import random
```

```
C:\Users\virgi\miniconda3\envs\reob_env\lib\site-packages\tqdm\auto.py:22: TqdmWarning: IPProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm
```

```
In [2]: 1 # Load a pretrained model yolov8x, yolov8x-seg(with segmentation), yolov8x-cls(with classification)
        2 model = YOLO("yolov8n-seg.pt")
```

### ▼ Detection

See [Detection Docs](#) for usage examples with these models.

Model	size (pixels)	mAP <sup>val</sup> 50-95	Speed CPU ONNX (ms)	Speed A100 TensorRT (ms)	params (M)	FLOPs (B)
<a href="#">YOLOv8n</a>	640	37.3	80.4	0.99	3.2	8.7
<a href="#">YOLOv8s</a>	640	44.9	128.4	1.20	11.2	28.6
<a href="#">YOLOv8m</a>	640	50.2	234.7	1.83	25.9	78.9
<a href="#">YOLOv8l</a>	640	52.9	375.2	2.39	43.7	165.2
<a href="#">YOLOv8x</a>	640	53.9	479.1	3.53	68.2	257.8

- mAP<sup>val</sup> values are for single-model single-scale on [COCO val2017](#) dataset.  
Reproduce by `yolo val detect data=coco.yaml device=0`
- Speed averaged over COCO val images using an [Amazon EC2 P4d](#) instance.  
Reproduce by `yolo val detect data=coco128.yaml batch=1 device=0/cpu`

### ▼ Segmentation

See [Segmentation Docs](#) for usage examples with these models.

Model	size (pixels)	mAP <sup>box</sup> 50-95	mAP <sup>mask</sup> 50-95	Speed CPU ONNX (ms)	Speed A100 TensorRT (ms)	params (M)	FLOPs (B)
<a href="#">YOLOv8n-seg</a>	640	36.7	30.5	96.1	1.21	3.4	12.6
<a href="#">YOLOv8s-seg</a>	640	44.6	36.8	155.7	1.47	11.8	42.6
<a href="#">YOLOv8m-seg</a>	640	49.9	40.8	317.0	2.18	27.3	110.2
<a href="#">YOLOv8l-seg</a>	640	52.3	42.6	572.4	2.79	46.0	220.5
<a href="#">YOLOv8x-seg</a>	640	53.4	43.4	712.1	4.02	71.8	344.1

- mAP<sup>val</sup> values are for single-model single-scale on [COCO val2017](#) dataset.  
Reproduce by `yolo val segment data=coco.yaml device=0`
- Speed averaged over COCO val images using an [Amazon EC2 P4d](#) instance.  
Reproduce by `yolo val segment data=coco128-seg.yaml batch=1 device=0/cpu`

### ▼ Classification

See [Classification Docs](#) for usage examples with these models.

Model	size (pixels)	acc top1	acc top5	Speed CPU ONNX (ms)	Speed A100 TensorRT (ms)	params (M)	FLOPs (B) at 640
<a href="#">YOLOv8n-cls</a>	224	66.6	87.0	12.9	0.31	2.7	4.3
<a href="#">YOLOv8s-cls</a>	224	72.3	91.1	23.4	0.35	6.4	13.5
<a href="#">YOLOv8m-cls</a>	224	76.4	93.2	85.4	0.62	17.0	42.7
<a href="#">YOLOv8l-cls</a>	224	78.0	94.1	163.0	0.87	37.5	99.7
<a href="#">YOLOv8x-cls</a>	224	78.4	94.3	232.0	1.01	57.4	154.8

- acc values are model accuracies on the [ImageNet](#) dataset validation set.  
Reproduce by `yolo val classify data=path/to/ImageNet device=0`
- Speed averaged over ImageNet val images using an [Amazon EC2 P4d](#) instance.  
Reproduce by `yolo val classify data=path/to/ImageNet batch=1 device=0/cpu`



# YOLO v8

## Inference

```
In [3]: 1 img = cv2.imread('bus.jpg') #"https://ultralytics.com/images/bus.jpg"
        2 img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
```

```
In [4]: 1 plt.imshow(img)
```


```
Out[4]: <matplotlib.image.AxesImage at 0x1b7f7d66d10>
```



```
In [5]: 1 r = model.predict(img)
        2 # r = model(img)
```

Ultralytics YOLOv8.0.39 Python-3.10.9 torch-1.13.1 CUDA:0 (NVIDIA GeForce RTX 2060, 6144MiB)  
YOLOv8n-seg summary (fused): 195 layers, 3404320 parameters, 0 gradients, 12.6 GFLOPs


0: 640x480 4 persons, 1 bus, 1 skateboard, 24.1ms





# YOLO v8

## Inference




```
In [6]: 1 names = model.names
2 colors = [tuple([random.randint(0, 255) for _ in range(3)]) for _ in range(len(names))]
3 roi = np.zeros(img.shape).astype(np.int32)
4 for box, mask in zip(r[0].boxes, r[0].masks):
5     # Bounding box & Category
6     bbox = box.xyxy.cpu().numpy()[0].astype(np.int32)
7     category = int(box.cls.cpu().numpy()[0])
8     img = cv2.rectangle(img, (bbox[0], bbox[1]), (bbox[2], bbox[3]), (colors[category]), 5)
9     label = names[category] + ' ' + str(box.conf.cpu().numpy()[0])
10    (w, h), _ = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX, 0.6, 1)
11    # Prints the text.
12    img = cv2.rectangle(img, (bbox[0], bbox[1] - 20), (bbox[0] + w, bbox[1]), colors[category], -1)
13    img = cv2.putText(img, label, (bbox[0], bbox[1] - 5), cv2.FONT_HERSHEY_SIMPLEX, 0.6, (0,0,0), 1)
14    # Mask
15    m = mask.data.cpu().numpy()
16    resized = cv2.resize(m, (img.shape[1], img.shape[0]), interpolation = cv2.INTER_AREA)
17    contours, _ = cv2.findContours(np.array(resized, np.uint8), cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
18    cnt = []
19    area = 0
20    for c in contours:
21        a = cv2.contourArea(c)
22        if a > area:
23            cnt = c
24            area = a
25    roi = cv2.drawContours(roi, [cnt], -1, colors[category], -1)
26    blended = ((0.6 * img) + (0.4 * roi)).astype("uint8")
27    plt.imshow(blended)
28    plt.show()
```

Prediction contains a list of bounding boxes and masks  
REMEMBER: it is a segmentation model



# YOLO v8

## Inference




```
In [6]: 1 names = model.names
2 colors = [tuple([random.randint(0, 255) for _ in range(3)]) for _ in range(len(names))]
3 roi = np.zeros(img.shape).astype(np.int32)
4 for box, mask in zip(r[0].boxes, r[0].masks):
5     # Bounding box & Category
6     bbox = box.xyxy.cpu().numpy()[0].astype(np.int32)
7     category = int(box.cls.cpu().numpy()[0])
8     img = cv2.rectangle(img, (bbox[0], bbox[1]), (bbox[2], bbox[3]), (colors[category]), 5)
9     label = names[category] + ' ' + str(box.conf.cpu().numpy()[0])
10    (w, h), _ = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX, 0.6, 1)
11    # Prints the text.
12    img = cv2.rectangle(img, (bbox[0], bbox[1] - 20), (bbox[0] + w, bbox[1]), colors[category], -1)
13    img = cv2.putText(img, label, (bbox[0], bbox[1] - 5), cv2.FONT_HERSHEY_SIMPLEX, 0.6, (0,0,0), 1)
14    # Mask
15    m = mask.data.cpu().numpy()
16    resized = cv2.resize(m, (img.shape[1], img.shape[0]), interpolation = cv2.INTER_AREA)
17    contours, _ = cv2.findContours(np.array(resized, np.uint8), cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
18    cnt = []
19    area = 0
20    for c in contours:
21        a = cv2.contourArea(c)
22        if a > area:
23            cnt = c
24            area = a
25    roi = cv2.drawContours(roi, [cnt], -1, colors[category], -1)
26    blended = ((0.6 * img) + (0.4 * roi)).astype("uint8")
27    plt.imshow(blended)
28    plt.show()
```

**Bounding box attributes (box)**  
xyxy: returns two extremes of the bounding box  
cls: class of the prediction  
conf: probability of the prediction



# YOLO v8

## Inference



```
In [6]: 1 names = model.names
2 colors = [tuple([random.randint(0, 255) for _ in range(3)]) for _ in range(len(names))]
3 roi = np.zeros(img.shape).astype(np.int32)
4 for box, mask in zip(r[0].boxes, r[0].masks):
5     # Bounding box & Category
6     bbox = box.xyxy.cpu().numpy()[0].astype(np.int32)
7     category = int(box.cls.cpu().numpy()[0])
8     img = cv2.rectangle(img, (bbox[0], bbox[1]), (bbox[2], bbox[3]), (colors[category]), 5)
9     label = names[category] + ' ' + str(box.conf.cpu().numpy()[0])
10    (w, h), _ = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX, 0.6, 1)
11    # Prints the text.
12    img = cv2.rectangle(img, (bbox[0], bbox[1] - 20), (bbox[0] + w, bbox[1]), colors[category], -1)
13    img = cv2.putText(img, label, (bbox[0], bbox[1] - 5), cv2.FONT_HERSHEY_SIMPLEX, 0.6, (0,0,0), 1)
14    # Mask
15    m = mask.data.cpu().numpy()
16    resized = cv2.resize(m, (img.shape[1], img.shape[0]), interpolat
17    contours, _ = cv2.findContours(np.array(resized, np.uint8), cv2.
18    cnt = []
19    area = 0
20    for c in contours:
21        a = cv2.contourArea(c)
22        if a > area:
23            cnt = c
24            area = a
25    roi = cv2.drawContours(roi, [cnt], -1, colors[category], -1)
26    blended = ((0.6 * img) + (0.4 * roi)).astype("uint8")
27    plt.imshow(blended)
28    plt.show()
```

**Bounding box attributes (box)**  
xyxy: returns two extremes of the bounding box  
cls: class of the prediction  
conf: probability of the prediction

### Bounding box attributes (box)


box.xyxy # box with xyxy format, (N, 4)  
box.xywh # box with xywh format, (N, 4)  
box.xyxy\_n # box with xyxy format but normalized, (N, 4)  
box.xywh\_n # box with xywh format but normalized, (N, 4)  
box.conf # confidence score, (N, 1)  
box.cls # cls, (N, 1)  
box.data # raw bboxes tensor, (N, 6) or boxes.boxes .





# YOLO v8

## Inference




```
In [6]: 1 names = model.names
2 colors = [tuple([random.randint(0, 255) for _ in range(3)]) for _ in range(len(names))]
3 roi = np.zeros(img.shape).astype(np.int32)
4 for box, mask in zip(r[0].boxes, r[0].masks):
5     # Bounding box & Category
6     bbox = box.xyxy.cpu().numpy()[0].astype(np.int32)
7     category = int(box.cls.cpu().numpy()[0])
8     img = cv2.rectangle(img, (bbox[0], bbox[1]), (bbox[2], bbox[3]), (colors[category]), 5)
9     label = names[category] + ' ' + str(box.conf.cpu().numpy()[0])
10    (w, h), _ = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX, 0.6, 1)
11    # Prints the text.
12    img = cv2.rectangle(img, (bbox[0], bbox[1] - 20), (bbox[0] + w, bbox[1]), colors[category], -1)
13    img = cv2.putText(img, label, (bbox[0], bbox[1] - 5), cv2.FONT_HERSHEY_SIMPLEX, 0.6, (0,0,0), 1)
14    # Mask
15    m = mask.data.cpu().numpy()
16    resized = cv2.resize(m, (img.shape[1], img.shape[0]), interpolation = cv2.INTER_AREA)
17    contours, _ = cv2.findContours(np.array(resized, np.uint8), cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
18    cnt = []
19    area = 0
20    for c in contours:
21        a = cv2.contourArea(c)
22        if a > area:
23            cnt = c
24            area = a
25    roi = cv2.drawContours(roi, [cnt], -1, colors[category], -1)
26    blended = ((0.6 * img) + (0.4 * roi)).astype("uint8")
27    plt.imshow(blended)
28    plt.show()
```

**Mask attributes (mask)**  
data: mask of the object



# YOLO v8



## Train

```
In [1]: 1 from ultralytics import YOLO
```

```
In [2]: 1 model = YOLO("yolov8n-seg.pt")
```

```
In [*]: 1 model.train(data="open_data/dataset.yaml", epochs=10, imgsz=640, name='test1') #image size to be divisible by 32
```

Ultralytics YOLOv8.0.39 Python-3.10.9 torch-1.13.1 CUDA:0 (NVIDIA GeForce RTX 2060, 6144MiB)

yolo\engine\trainer: task=segment, mode=train, model=yolov8n-seg.pt, data=open\_data/dataset.yaml, epochs=10, patience=50, bat


Key	Value	Description
model	None	path to model file, i.e. yolov8n.pt, yolov8n.yaml
data	None	path to data file, i.e. coco128.yaml
epochs	100	number of epochs to train for
batch	16	number of images per batch (-1 for AutoBatch)
imgsz	640	size of input images as integer or w,h
save	True	save train checkpoints and predict results
save_period	-1	Save checkpoint every x epochs (disabled if < 1)
cache	False	True/ram, disk or False. Use cache for data loading
device	None	device to run on, i.e. cuda device=0 or device=0,1,2,3 or device=cpu
workers	8	number of worker threads for data loading (per RANK if DDP)
name	None	experiment name
pretrained	False	whether to use a pretrained model
optimizer	'SGD'	optimizer to use, choices=['SGD', 'Adam', 'AdamW', 'RMSProp']
verbose	False	whether to print verbose output
lr0	0.01	initial learning rate (i.e. SGD=1E-2, Adam=1E-3)
lrf	0.01	final learning rate (lr0 * lrf)

project=None, name=test1, exist\_ok=False, p  
s=False, image\_weights=False, rect=False,  
ask\_ratio=4, dropout=False, val=True, spli  
alse, dnn=False, plots=True, source=ultral  
labels=False, hide\_conf=False, vid\_stride=  
ne, retina\_masks=False, boxes=True, format  
e, opset=None, workspace=4, nms=False, lr0  
ntum=0.8, warmup\_bias\_lr=0.1, box=7.5, cls  
sv\_v=0.4, degrees=0.0, translate=0.1, scal  
copy\_paste=0.0, cfg=None, v5loader=False, t

arguments  
[3, 16, 3, 2]  
[16, 32, 3, 2]  
[32, 32, 1, True]



# YOLO v8



## Train

```
In [1]: 1 from ultralytics import YOLO
```

```
In [2]: 1 model = YOLO("yolov8n-seg.pt")
```

```
In [*]: 1 model.train(data="open_data/dataset.yaml", epochs=10, imgsz=640, name='test1') #image size to be divisible by 32
```

Ultralytics YOLOv8.0.39 Python-3.10.9 torch-1.13.1 CUDA:0 (NVIDIA GeForce RTX 2060, 6144MiB)

yolo\engine\trainer: task=segment, mode=train, model=yolov8n-seg.pt, data=open\_data/dataset.yaml, epochs=10, patience=50, batch=16, imgsz=640, save=True, save\_period=-1, cache=False, device=None, workers=8, project=None, name=test1, exist\_ok=False, pretrained=False, optimizer=SGD, verbose=True, seed=0, deterministic=True, single\_cls=False, image\_weights=False, rect=False,

True, mask\_ratio=4, dropout=False, val=True, split\_val\_half=False, dnn=False, plots=True, source=ultralytics, hide\_labels=False, hide\_conf=False, vid\_stride=1, classes=None, retina\_masks=False, boxes=True, format=torch, opset=None, workspace=4, nms=False, lr0=0.01, momentum=0.8, warmup\_bias\_lr=0.1, box=7.5, cls=0.7, hsv\_v=0.4, degrees=0.0, translate=0.1, scale=0.0, copy\_paste=0.0, cfg=None, v5loader=False, t


arguments  
[3, 16, 3, 2]  
[16, 32, 3, 2]  
[32, 32, 1, True]

Key	Value	Description
momentum	0.937	SGD momentum/Adam beta1
weight_decay	0.0005	optimizer weight decay 5e-4
warmup_epochs	3.0	warmup epochs (fractions ok)
warmup_momentum	0.8	warmup initial momentum
warmup_bias_lr	0.1	warmup initial bias lr
box	7.5	box loss gain
cls	0.5	cls loss gain (scale with pixels)
dfi	1.5	dfi loss gain
fl_gamma	0.0	focal loss gamma (efficientDet default gamma=1.5)
label_smoothing	0.0	label smoothing (fraction)
nbs	64	nominal batch size
overlap_mask	True	masks should overlap during training (segment train only)
mask_ratio	4	mask downsample ratio (segment train only)
dropout	0.0	use dropout regularization (classify train only)
val	True	validate/test during training



# YOLO v8

## Train



runs



segment



test1 (name)



weights



args.yaml



BoxF1\_curve.png



BoxP\_curve.png



BoxPR\_curve.png



BoxR\_curve.png



confusion\_matrix.png



events.out.tfevents.1679478250.LAPTOP-K9C286D7.16568.0



MaskF1\_curve.png



MaskP\_curve.png




MaskPR\_curve.png



MaskR\_curve.png



results.csv



results.png



train\_batch0.jpg




train\_batch1.jpg



train\_batch2.jpg



val\_batch0\_labels.jpg



val\_batch0\_pred.jpg



val\_batch1\_labels.jpg



val\_batch1\_pred.jpg




val\_batch2\_labels.jpg



val\_batch2\_pred.jpg



## Export



```
Ultralytics YOLOv8.0.39 Python-3.10.9 torch-1.13.1 CPU
YOLOv8n-seg summary (fused): 195 layers, 3258259 parameters, 0 gradients, 12.0 GFLOPs
```

```
requirements: YOLOv8 requirement "onnx>=1.12.0" not found, attempting AutoUpdate...
```

Requirement already satisfied: protobuf<4,>=3.20.2 in c:\users\virgi\miniconda3\envs\reob\_env\lib\site-packages (from onnx>=1.12.0) (3.20.3)

```
Requirement already satisfied: typing-extensions>=3.6.2.1 in c:\users\virgi\miniconda3\envs\reob_env\lib\site-packages (from onnx>=1.12.0) (4.4.0)
```

Requirement already satisfied: numpy>=1.16.6 in c:\users\virgi\miniconda3\envs\reob\_env\lib\site-packages (from onnx>=1.12.0 (1.23.5))

```
requirements: 1 package updated per ['onnx>=1.1
```

requirements: Restart runtime or rerun command

```
ONNX: starting export with onnx 1.13.1...
```

```
ONNX: export success 4.3s, saved as runs\segme
```

Export complete (4.8s)

Results saved to G:\Mi unidad\DOCENCIA\MARSI\re

```
Predict:      yolo task=segment mode=predict
```

```
Validate:      yolo task=segment mode=val mod
```

Visualize: <https://netron.app>

```
'runs\\segment\\test1\\weights\\best.onnx'
```


Key	Value	Description
format	'torchscript'	format to export to
imgsz	640	image size as scalar or (h, w) list, i.e. (640, 480)
keras	False	use Keras for TF SavedModel export
optimize	False	TorchScript: optimize for mobile
half	False	FP16 quantization
int8	False	INT8 quantization
dynamic	False	ONNX/TF/TensorRT: dynamic axes
simplify	False	ONNX: simplify model
opset	None	ONNX: opset version (optional, defaults to latest)
workspace	4	TensorRT: workspace size (GB)
nms	False	CoreML: add NMS





# YOLO v8

## Export



Format	format Argument	Model	Metadata
PyTorch	-	yolov8n.pt	✓
TorchScript	torchscript	yolov8n.torchscript	✓
ONNX	onnx	yolov8n.onnx	✓
OpenVINO	openvino	yolov8n_openvino_model/	✓
TensorRT	engine	yolov8n.engine	✓
CoreML	coreml	yolov8n.mlmodel	✓
TF SavedModel	saved_model	yolov8n_saved_model/	✓
TF GraphDef	pb	yolov8n.pb	✗
TF Lite	tflite	yolov8n.tflite	✓
TF Edge TPU	edgetpu	yolov8n_edgetpu.tflite	✓
TF.js	tfjs	yolov8n_web_model/	✓
PaddlePaddle	paddle	yolov8n_paddle_model/	✓

